

The Tale of the PCP Theorem

The PCP Theorem and Hardness of Approximation

Jatin Arora Shaan Vaidya

Computational Complexity

Outline

- 1 Limits of Approximation
- 2 GAP problems
- 3 PCP theorem
- 4 Proof of Hardness of Approximation
- 5 Alternate View of PCP theorem

Approximate solutions to NP-hard optimization problems

- MAX-3SAT is NP-hard
- Can't hope for a fast algorithm that always gets the best solution
- Can we at least guarantee to get a “pretty good” solution efficiently?

Definition

Approximation of MAX-3SAT:

For every $\rho \leq 1$, an algorithm A is a ρ – approximation algorithm for MAX-3SAT if for every 3CNF formula φ with m clauses, $A(\varphi)$ outputs an assignment satisfying at least $\rho \cdot \text{val}(\varphi)$ of φ 's clauses.

- Greedy algorithm - $1/2$ -approximation
- Best known - $7/8$ approximation algorithm

Definition

Approximation of MAX-3SAT:

For every $\rho \leq 1$, an algorithm A is a ρ – approximation algorithm for MAX-3SAT if for every 3CNF formula φ with m clauses, $A(\varphi)$ outputs an assignment satisfying at least $\rho \cdot \text{val}(\varphi)$ of φ 's clauses.

- Greedy algorithm - $1/2$ -approximation
- Best known - $7/8$ approximation algorithm
- Can we do better? Or is there a limit to approximation?

- How does one study the hardness of approximation problems?

- MAXCLIQUE: Given a graph G , output the vertices in its largest clique
- SET COVERING PROBLEM: Given a collection of sets S_1, S_2, \dots, S_m that cover a universe $U = \{1, 2, \dots, n\}$, find the smallest sub-collection of sets S_{i_1}, S_{i_2}, \dots that also cover the universe

- MAXCLIQUE: Given a graph G , output the vertices in its largest clique
- CLIQUE = $\{(G, k) \mid \text{graph } G \text{ has a clique of size } \geq k\}$
- SET COVERING PROBLEM: Given a collection of sets S_1, S_2, \dots, S_m that cover a universe $U = \{1, 2, \dots, n\}$, find the smallest sub-collection of sets S_{i_1}, S_{i_2}, \dots that also cover the universe
- SET-COVER = $\{(U; \{S_1, S_2, \dots, S_m\}, k) \mid \exists i_1 \leq i_2 \dots, i_k \leq m \text{ st } \bigcup_{j=1}^m S_{i_j} = U\}$

GAP problems

- MAXCLIQUE: Given a graph G , output the vertices in its largest clique
- CLIQUE = $\{(G, k) \mid \text{graph } G \text{ has a clique of size } \geq k\}$
- SET COVERING PROBLEM: Given a collection of sets S_1, S_2, \dots, S_m that cover a universe $U = \{1, 2, \dots, n\}$, find the smallest sub-collection of sets S_{i_1}, S_{i_2}, \dots that also cover the universe
- SET-COVER = $\{(U; \{S_1, S_2, \dots, S_m\}, k) \mid \exists i_1 \leq i_2 \leq \dots \leq i_k \leq m \text{ st } \bigcup_{j=1}^k S_{i_j} = U\}$

As in the case of computational problems, it would be nice if we could capture the hardness of the approximation problems via decision problems.

Definition

Promise problems

A promise problem $\Pi \subseteq \Sigma^*$ is specified by a pair of sets (YES, NO) such that $\text{YES}, \text{NO} \subseteq \Sigma^*$ and $\text{YES} \cap \text{NO} = \emptyset$

Definition

Promise problems

A promise problem $\Pi \subseteq \Sigma^*$ is specified by a pair of sets (YES, NO) such that $\text{YES}, \text{NO} \subseteq \Sigma^*$ and $\text{YES} \cap \text{NO} = \emptyset$

Gap problems are promise problems

The gap problem corresponding to the α -approximating MAX3SAT, called $\text{gap}_\alpha\text{-MAX3SAT}$ (for $\alpha \leq 1$)

The gap problem corresponding to the α -approximating MAX3SAT, called $\text{gap}_\alpha\text{-MAX3SAT}$ (for $\alpha \leq 1$)

$\text{gap}_\alpha\text{-MAX3SAT}$ is a promise problem whose (YES, NO) are as follows:

YES = $\{\langle \varphi, k \rangle \mid \text{there is an assignment satisfying } \geq k \text{ clauses of } \varphi\}$

NO = $\{\langle \varphi, k \rangle \mid \text{every assignment satisfies } < \alpha k \text{ clauses of } \varphi\}$

Equivalence of GAP problems and Approximation problems

For any $0 < \alpha < 1$, α -approximating MAX3SAT is polynomially equivalent to solving gap_α -MAX3SAT

Equivalence of GAP problems and Approximation problems

Proof. (\Rightarrow)

Suppose there is an α -approximation algorithm **A** to MAX3SAT.

Consider the following algorithm **B** for $\text{gap}_\alpha\text{-MAX3SAT}$

B : On input $\langle \varphi, k \rangle$

1. Run **A** on φ and let $k' = \mathbf{A}(\varphi)$
2. Accept iff $k' \geq \alpha k$.

Equivalence of GAP problems and Approximation problems

Proof. (\Rightarrow)

Suppose there is an α -approximation algorithm **A** to MAX3SAT.

Consider the following algorithm **B** for $\text{gap}_\alpha\text{-MAX3SAT}$

B : On input $\langle \varphi, k \rangle$

1. Run **A** on φ and let $k' = \mathbf{A}(\varphi)$
2. Accept iff $k' \geq \alpha k$.

Equivalence of GAP problems and Approximation problems

Proof. (\Rightarrow)

Suppose there is an α -approximation algorithm **A** to MAX3SAT.

Consider the following algorithm **B** for $\text{gap}_\alpha\text{-MAX3SAT}$

B : On input $\langle \varphi, k \rangle$

1. Run **A** on φ and let $k' = \mathbf{A}(\varphi)$
2. Accept iff $k' \geq \alpha k$.

(\Leftarrow) Suppose instead there is an algorithm **B** that solves $\text{gap}_\alpha\text{-MAX3SAT}$

A : On input φ

1. Let m be the number of clauses of φ
2. Run **B** on $\langle \varphi, 1 \rangle, \langle \varphi, 2 \rangle, \langle \varphi, 3 \rangle, \dots, \langle \varphi, m \rangle$.
3. Let the largest k such that **B** accepted $\langle \varphi, k \rangle$
4. Output αk

Equivalence of GAP problems and Approximation problems

Thus, to show that approximating MAX3SAT is hard, it suffices (and is necessary) to show that $\text{gap}_\alpha\text{-MAX3SAT}$ is hard.

Interactive Proof Systems

- A proof system consists of a verifier V and prover P

Interactive Proof Systems

- A proof system consists of a verifier V and prover P
- Given a statement x , such as φ is satisfiable P produces a candidate proof π for the statement φ

Interactive Proof Systems

- A proof system consists of a verifier V and prover P
- Given a statement x , such as φ is satisfiable P produces a candidate proof π for the statement φ
- The verifier V then reads the statement-proof pair (φ, π) and either accepts or rejects the proof π for φ .

Interactive Proof Systems

- A proof system consists of a verifier V and prover P
- Given a statement x , such as φ is satisfiable P produces a candidate proof π for the statement φ
- The verifier V then reads the statement-proof pair (φ, π) and either accepts or rejects the proof π for φ .
- Completeness and Soundness properties define many complexity classes of importance

Probabilistically Checkable Proof Systems

Definition

PCP verifier

Let L be a language and $q, r : \mathbb{N} \rightarrow \mathbb{N}$. We say that L has an $(r(n), q(n))$ -PCP verifier if there is a polynomial-time probabilistic algorithm \mathcal{V} satisfying:

Probabilistically Checkable Proof Systems

Definition

PCP verifier

Let L be a language and $q, r : \mathbb{N} \rightarrow \mathbb{N}$. We say that L has an $(r(n), q(n))$ -PCP verifier if there is a polynomial-time probabilistic algorithm \mathcal{V} satisfying:

Efficiency On input $x \in \{0, 1\}^n$ and given random access to the proof $\pi \in \{0, 1\}^*$, \mathcal{V} uses at most $r(n)$ random coins and makes at most $q(n)$ queries to locations of π to decide x .

Probabilistically Checkable Proof Systems

Definition

PCP verifier

Let L be a language and $q, r : \mathbb{N} \rightarrow \mathbb{N}$. We say that L has an $(r(n), q(n))$ -PCP verifier if there is a polynomial-time probabilistic algorithm \mathcal{V} satisfying:

Efficiency On input $x \in \{0, 1\}^n$ and given random access to the proof $\pi \in \{0, 1\}^*$, \mathcal{V} uses at most $r(n)$ random coins and makes at most $q(n)$ queries to locations of π to decide x .

Completeness $x \in L \implies \exists \pi \in \{0, 1\}^*, \Pr[\mathcal{V}^\pi(x) = 1] = 1$

Probabilistically Checkable Proof Systems

Definition

PCP verifier

Let L be a language and $q, r : \mathbb{N} \rightarrow \mathbb{N}$. We say that L has an $(r(n), q(n))$ -PCP verifier if there is a polynomial-time probabilistic algorithm \mathcal{V} satisfying:

Efficiency On input $x \in \{0, 1\}^n$ and given random access to the proof $\pi \in \{0, 1\}^*$, \mathcal{V} uses at most $r(n)$ random coins and makes at most $q(n)$ queries to locations of π to decide x .

Completeness $x \in L \implies \exists \pi \in \{0, 1\}^*, \Pr[\mathcal{V}^\pi(x) = 1] = 1$

Soundness $x \notin L \implies \forall \pi \Pr[\mathcal{V}^\pi(x) = 1] \leq 1/2$

Probabilistically Checkable Proof Systems

Definition

PCP verifier

Let L be a language and $q, r : \mathbb{N} \rightarrow \mathbb{N}$. We say that L has an $(r(n), q(n))$ -PCP verifier if there is a polynomial-time probabilistic algorithm \mathcal{V} satisfying:

Efficiency On input $x \in \{0, 1\}^n$ and given random access to the proof $\pi \in \{0, 1\}^*$, \mathcal{V} uses at most $r(n)$ random coins and makes at most $q(n)$ queries to locations of π to decide x .

Completeness $x \in L \implies \exists \pi \in \{0, 1\}^*, \Pr[\mathcal{V}^\pi(x) = 1] = 1$

Soundness $x \notin L \implies \forall \pi \Pr[\mathcal{V}^\pi(x) = 1] \leq 1/2$

L is in $\text{PCP}(r(n), q(n))$ if L has a $(O(r(n)), O(q(n)))$ -PCP verifier

Some remarks

- The verifier could be either non-adaptive or adaptive (i.e., the locations probed by the verifier could depend on the earlier probes)

Some remarks

- The verifier could be either non-adaptive or adaptive (i.e., the locations probed by the verifier could depend on the earlier probes)
- If the verifier is non-adaptive then the size of the proof is bounded above by $q(n) \cdot 2^{r(n)}$ while if the verifier is adaptive, it is bounded by $2^{r(n)+q(n)}$
- Assume non-adaptive unless otherwise specified

Some remarks

- The verifier could be either non-adaptive or adaptive (i.e., the locations probed by the verifier could depend on the earlier probes)
- If the verifier is non-adaptive then the size of the proof is bounded above by $q(n) \cdot 2^{r(n)}$ while if the verifier is adaptive, it is bounded by $2^{r(n)+q(n)}$
- Assume non-adaptive unless otherwise specified
- $\text{PCP}(r(n), q(n)) \subseteq \text{NTIME}(2^{r(n)} \cdot q(n))$

Some remarks

- The verifier could be either non-adaptive or adaptive (i.e., the locations probed by the verifier could depend on the earlier probes)
- If the verifier is non-adaptive then the size of the proof is bounded above by $q(n) \cdot 2^{r(n)}$ while if the verifier is adaptive, it is bounded by $2^{r(n)+q(n)}$
- Assume non-adaptive unless otherwise specified
- $\text{PCP}(r(n), q(n)) \subseteq \text{NTIME}(2^{r(n)} \cdot q(n))$
- It follows from definition that:
 - $\text{NP} = \text{PCP}_{1,0}(0, \text{poly}(n))$
 - $\text{BPP} = \text{PCP}_{2/3,1/3}(\text{poly}(n), 0)$
 - $\text{P} = \text{PCP}_{1,0}(0, 0)$.

Error reduction in PCP

A PCP verifier with soundness $1/2$ that uses r coins and makes q queries can be converted into a PCP verifier using $c \cdot r$ coins and $c \cdot q$ queries with soundness $1/2^c$ by just repeating its execution c times

The PCP Theorem

PCP Theorem

$$\text{NP} = \text{PCP}(\log n, 1)$$

Hardness of Approximation Using PCP Theorem

Consider the generalization of the MAX-3SAT problem and its corresponding GAP problem

Constraint satisfaction problems (CSP)

A q CSP instance φ is a collection of boolean constraints $\varphi_1, \dots, \varphi_m$ such that each φ_i depends on at most q of the input variables

Let $\text{val}(\varphi)$ denote the maximum fraction of constraints that can be satisfied by any assignment

Hardness of Approximation Using PCP Theorem

Consider the generalization of the MAX-3SAT problem and its corresponding GAP problem

Gap CSP

$\text{gap}_\alpha\text{-}q\text{CSP}$ is the promise problem whose (YES, NO) are as follows:

$$\text{YES} = \{\varphi \mid \text{val}(\varphi) = 1\}$$

$$\text{NO} = \{\varphi \mid \text{val}(\varphi) < \alpha\}$$

gap_α - q CSP is NP-hard

Theorem

$\exists q \in \mathbb{N}, \alpha \in (0, 1) : \text{gap}_\alpha$ - q CSP is NP-hard

$\text{gap}_\alpha\text{-}q\text{CSP}$ is NP-hard

Proof

- $3\text{SAT} \in \text{NP} \implies 3\text{SAT} \in \text{PCP}(\log n, 1)$

$\text{gap}_\alpha\text{-}q\text{CSP}$ is NP-hard

Proof

- $3\text{SAT} \in \text{NP} \implies 3\text{SAT} \in \text{PCP}(\log n, 1)$
- Let \mathcal{V} be the $\text{PCP}(c \cdot \log n, d)$ verifier for 3SAT

Proof

- $3\text{SAT} \in \text{NP} \implies 3\text{SAT} \in \text{PCP}(\log n, 1)$
- Let \mathcal{V} be the $\text{PCP}(c \cdot \log n, d)$ verifier for 3SAT
- Let $\mathcal{V}_{x,r}$ be a boolean constraint which evaluates to 1 iff \mathcal{V} accepts π , for input 3SAT instance x and r coins

$\text{gap}_\alpha\text{-}q\text{CSP}$ is NP-hard

Proof

- $3\text{SAT} \in \text{NP} \implies 3\text{SAT} \in \text{PCP}(\log n, 1)$
- Let \mathcal{V} be the $\text{PCP}(c \cdot \log n, d)$ verifier for 3SAT
- Let $\mathcal{V}_{x,r}$ be a boolean constraint which evaluates to 1 iff \mathcal{V} accepts π , for input 3SAT instance x and r coins
- $\mathcal{V}_{x,r}$ is a formula on d variables which correspond to the queries

$\text{gap}_\alpha\text{-}q\text{CSP}$ is NP-hard

Proof

- $3\text{SAT} \in \text{NP} \implies 3\text{SAT} \in \text{PCP}(\log n, 1)$
- Let \mathcal{V} be the $\text{PCP}(c \cdot \log n, d)$ verifier for 3SAT
- Let $\mathcal{V}_{x,r}$ be a boolean constraint which evaluates to 1 iff \mathcal{V} accepts π , for input 3SAT instance x and r coins
- $\mathcal{V}_{x,r}$ is a formula on d variables which correspond to the queries
- $\mathcal{V}_{x,r}$ is poly-sized since \mathcal{V} is a polynomial verifier

$\text{gap}_\alpha\text{-}q\text{CSP}$ is NP-hard

Proof

- $3\text{SAT} \in \text{NP} \implies 3\text{SAT} \in \text{PCP}(\log n, 1)$
- Let \mathcal{V} be the $\text{PCP}(c \cdot \log n, d)$ verifier for 3SAT
- Let $\mathcal{V}_{x,r}$ be a boolean constraint which evaluates to 1 iff \mathcal{V} accepts π , for input 3SAT instance x and r coins
- $\mathcal{V}_{x,r}$ is a formula on d variables which correspond to the queries
- $\mathcal{V}_{x,r}$ is poly-sized since \mathcal{V} is a polynomial verifier
- So for any x , we construct the $d\text{CSP}$ instance $\varphi = \{\mathcal{V}_{x,r}\}_{r \in \{0,1\}^{c \cdot \log n}}$

Hardness of Approximation Using PCP Theorem

Proof

- Since $r \in \{0,1\}^{c \cdot \log n}$, φ can be constructed in polynomial time

Hardness of Approximation Using PCP Theorem

Proof

- Since $r \in \{0,1\}^{c \cdot \log n}$, φ can be constructed in polynomial time
- $x \in 3\text{SAT} \implies \text{val}(\varphi) = 1$

Hardness of Approximation Using PCP Theorem

Proof

- Since $r \in \{0,1\}^{c \cdot \log n}$, φ can be constructed in polynomial time
- $x \in 3\text{SAT} \implies \text{val}(\varphi) = 1$
- $x \notin 3\text{SAT} \implies \text{val}(\varphi) \leq 1/2$

Hardness of Approximation Using PCP Theorem

Proof

- Since $r \in \{0, 1\}^{c \cdot \log n}$, φ can be constructed in polynomial time
- $x \in 3\text{SAT} \implies \text{val}(\varphi) = 1$
- $x \notin 3\text{SAT} \implies \text{val}(\varphi) \leq 1/2$
- Hence, φ is a valid instance for the promise problem $\text{gap}_{1/2}\text{-}d\text{CSP}$

Hardness of Approximation Using PCP Theorem

Proof

- Since $r \in \{0, 1\}^{c \cdot \log n}$, φ can be constructed in polynomial time
- $x \in 3\text{SAT} \implies \text{val}(\varphi) = 1$
- $x \notin 3\text{SAT} \implies \text{val}(\varphi) \leq 1/2$
- Hence, φ is a valid instance for the promise problem $\text{gap}_{1/2}\text{-}d\text{CSP}$
- $3\text{SAT} \leq_P \text{gap}_{1/2}\text{-}d\text{CSP}$

PCP Theorem Using Hardness of Approximation

We saw that PCP Theorem \implies Gap problem is hard

PCP Theorem Using Hardness of Approximation

We saw that PCP Theorem \implies Gap problem is hard

What about the converse?

PCP Theorem Using Hardness of Approximation

We saw that PCP Theorem \implies Gap problem is hard

What about the converse?

It turns out the converse is also true i.e. hardness of the Gap problem is an alternate view of the PCP theorem

PCP Theorem Using Hardness of Approximation

- We have that $\text{gap}_{1/2-q}\text{CSP}$ is NP-hard for some q

PCP Theorem Using Hardness of Approximation

- We have that $\text{gap}_{1/2-q}\text{CSP}$ is NP-hard for some q
- PCP Theorem: $\text{NP} = \text{PCP}(\log n, 1)$

PCP Theorem Using Hardness of Approximation

- We have that $\text{gap}_{1/2-q}\text{CSP}$ is NP-hard for some q
- PCP Theorem: $\text{NP} = \text{PCP}(\log n, 1)$
- $\text{PCP}(\log n, 1) \subseteq \text{NP}$

PCP Theorem Using Hardness of Approximation

- We have that $\text{gap}_{1/2-q}\text{CSP}$ is NP-hard for some q
- PCP Theorem: $\text{NP} = \text{PCP}(\log n, 1)$
- $\text{PCP}(\log n, 1) \subseteq \text{NP}$
- If we show that $\text{NP} \subseteq \text{PCP}(\log n, 1)$, we are done

PCP Theorem Using Hardness of Approximation

- $L \in NP \implies L \leq_P \text{gap}_{1/2}\text{-}q\text{CSP}$

PCP Theorem Using Hardness of Approximation

- $L \in NP \implies L \leq_P \text{gap}_{1/2}\text{-}q\text{CSP}$
- for any input $x \in L$, we construct φ_x in polynomial time

PCP Theorem Using Hardness of Approximation

- $L \in NP \implies L \leq_P \text{gap}_{1/2}\text{-}q\text{CSP}$
- for any input $x \in L$, we construct φ_x in polynomial time
- The PCP Verifier(\mathcal{V}) for φ_x expects a satisfying assignment to the variables from the prover

PCP Theorem Using Hardness of Approximation

- $L \in NP \implies L \leq_P \text{gap}_{1/2}\text{-}q\text{CSP}$
- for any input $x \in L$, we construct φ_x in polynomial time
- The PCP Verifier(\mathcal{V}) for φ_x expects a satisfying assignment to the variables from the prover
- It then picks a constraint φ_i , at random from φ_x

PCP Theorem Using Hardness of Approximation

- $L \in NP \implies L \leq_P \text{gap}_{1/2}\text{-}q\text{CSP}$
- for any input $x \in L$, we construct φ_x in polynomial time
- The PCP Verifier(\mathcal{V}) for φ_x expects a satisfying assignment to the variables from the prover
- It then picks a constraint φ_i , at random from φ_x
- This can be done with $O(\log n)$ random bits since φ_x is poly-sized

PCP Theorem Using Hardness of Approximation

- $L \in NP \implies L \leq_P \text{gap}_{1/2}\text{-}q\text{CSP}$
- for any input $x \in L$, we construct φ_x in polynomial time
- The PCP Verifier(\mathcal{V}) for φ_x expects a satisfying assignment to the variables from the prover
- It then picks a constraint φ_i , at random from φ_x
- This can be done with $O(\log n)$ random bits since φ_x is poly-sized
- Then it queries the assignment of q variables associated with φ_i

PCP Theorem Using Hardness of Approximation

- $L \in NP \implies L \leq_P \text{gap}_{1/2}\text{-}q\text{CSP}$
- for any input $x \in L$, we construct φ_x in polynomial time
- The PCP Verifier(\mathcal{V}) for φ_x expects a satisfying assignment to the variables from the prover
- It then picks a constraint φ_i , at random from φ_x
- This can be done with $O(\log n)$ random bits since φ_x is poly-sized
- Then it queries the assignment of q variables associated with φ_i
- \mathcal{V} accepts iff φ_i is satisfied

PCP Theorem Using Hardness of Approximation

- $x \in L \implies \text{val}(\varphi) = 1$, hence \mathcal{V} will accept with probability 1

PCP Theorem Using Hardness of Approximation

- $x \in L \implies \text{val}(\varphi) = 1$, hence \mathcal{V} will accept with probability 1
- $x \notin L \implies \text{val}(\varphi) \leq 1/2$, hence \mathcal{V} will accept with probability $\leq 1/2$

Two views of the PCP theorem

Proof View	Hardness of Approximation View
PCP Verifier (\mathcal{V})	CSP Instance (φ)
PCP proof (π)	Assignment to variables (\mathbf{u})
Length of proof	Number of variables (n)
Number of queries (q)	Arity of constraints (q)
Number of random bits (r)	Logarithm of number of constraints ($\log m$)
Soundness parameter (typically $1/2$)	Maximum of $\text{val}(\varphi)$ for a NO instance
$\text{NP} \subseteq \text{PCP}(\log n, 1)$	$\text{gap}_{1/2-q}\text{CSP}$ is NP-hard

Consider a slightly different definition of the PCP classes

Definition

$\text{PCP}_{c,s}^{\Sigma}(r, q)$ is the class of languages that have restricted verifiers that use r random bits, q queries to the proof $\pi \in \Sigma^*$ with

Completeness $x \in L \implies \exists \pi, \Pr[\mathcal{V}^{\pi}(x) = 1] \geq c$

Soundness $x \notin L \implies \forall \pi, \Pr[\mathcal{V}^{\pi}(x) = 1] \leq s$

Theorem

$$PCP_{c,1-\epsilon}^{\Sigma}[r, q] \subseteq PCP_{c,1-\epsilon/q}^{\Sigma}[r + \log q, 2]$$

Proof.

- $L \in PCP_{c,1-\epsilon}^{\Sigma}[r, q] \implies L$ has a (r, q) -verifier \mathcal{V}
- for $x \in L$, $\exists \pi$ such that \mathcal{V} accepts with probability $\geq c$
- for $x \notin L$, $\forall \pi$, \mathcal{V} accepts with probability $\leq s$
- Define $Perm(\pi) : [m]^q \rightarrow \Sigma^q$ st
 $Perm(\pi)(i_1, \dots, i_q) = \pi(i_1)\pi(i_2) \dots \pi(i_q)$

- Consider a verifier \mathcal{V}' that expects $\langle Perm(\pi), \pi \rangle$ as proof

2-query PCP

- Consider a verifier \mathcal{V}' that expects $\langle Perm(\pi), \pi \rangle$ as proof
- The prover actually sends $\langle \pi_1, \pi_2 \rangle$

2-query PCP

- Consider a verifier \mathcal{V}' that expects $\langle Perm(\pi), \pi \rangle$ as proof
- The prover actually sends $\langle \pi_1, \pi_2 \rangle$
- \mathcal{V}' then guesses the indices of queries (i_1, \dots, i_q) as \mathcal{V} would using r coins

2-query PCP

- Consider a verifier \mathcal{V}' that expects $\langle \text{Perm}(\pi), \pi \rangle$ as proof
- The prover actually sends $\langle \pi_1, \pi_2 \rangle$
- \mathcal{V}' then guesses the indices of queries (i_1, \dots, i_q) as \mathcal{V} would using r coins
- \mathcal{V}' queries $\pi_1(i_1, \dots, i_q)$ and proceeds as \mathcal{V} would

2-query PCP

- Consider a verifier \mathcal{V}' that expects $\langle \text{Perm}(\pi), \pi \rangle$ as proof
- The prover actually sends $\langle \pi_1, \pi_2 \rangle$
- \mathcal{V}' then guesses the indices of queries (i_1, \dots, i_q) as \mathcal{V} would using r coins
- \mathcal{V}' queries $\pi_1(i_1, \dots, i_q)$ and proceeds as \mathcal{V} would
- If \mathcal{V} would have rejected then \mathcal{V}' will reject

- Consider a verifier \mathcal{V}' that expects $\langle \text{Perm}(\pi), \pi \rangle$ as proof
- The prover actually sends $\langle \pi_1, \pi_2 \rangle$
- \mathcal{V}' then guesses the indices of queries (i_1, \dots, i_q) as \mathcal{V} would using r coins
- \mathcal{V}' queries $\pi_1(i_1, \dots, i_q)$ and proceeds as \mathcal{V} would
- If \mathcal{V} would have rejected then \mathcal{V}' will reject
- Otherwise, \mathcal{V}' selects $j \in [q]$ using $\log q$ coins

- Consider a verifier \mathcal{V}' that expects $\langle \text{Perm}(\pi), \pi \rangle$ as proof
- The prover actually sends $\langle \pi_1, \pi_2 \rangle$
- \mathcal{V}' then guesses the indices of queries (i_1, \dots, i_q) as \mathcal{V} would using r coins
- \mathcal{V}' queries $\pi_1(i_1, \dots, i_q)$ and proceeds as \mathcal{V} would
- If \mathcal{V} would have rejected then \mathcal{V}' will reject
- Otherwise, \mathcal{V}' selects $j \in [q]$ using $\log q$ coins
- \mathcal{V}' accepts iff j th symbol of $\pi_1(i_1, \dots, i_q) = \pi_2(i_j)$

$$x \in L \Rightarrow$$

2-query PCP

$x \in L \Rightarrow$

- $\exists \pi$ that \mathcal{V} accepts with probability $\geq c$

2-query PCP

$x \in L \Rightarrow$

- $\exists \pi$ that \mathcal{V} accepts with probability $\geq c$
- $\langle \text{Perm}(\pi), \pi \rangle$ is a proof that \mathcal{V}' accepts with probability $\geq c$

2-query PCP

$x \in L \Rightarrow$

- $\exists \pi$ that \mathcal{V} accepts with probability $\geq c$
- $\langle \text{Perm}(\pi), \pi \rangle$ is a proof that \mathcal{V}' accepts with probability $\geq c$

$x \notin L \Rightarrow$

2-query PCP

$x \in L \Rightarrow$

- $\exists \pi$ that \mathcal{V} accepts with probability $\geq c$
- $\langle \text{Perm}(\pi), \pi \rangle$ is a proof that \mathcal{V}' accepts with probability $\geq c$

$x \notin L \Rightarrow$

- For any π , \mathcal{V} rejects with probability $\geq \epsilon$

2-query PCP

$x \in L \Rightarrow$

- $\exists \pi$ that \mathcal{V} accepts with probability $\geq c$
- $\langle \text{Perm}(\pi), \pi \rangle$ is a proof that \mathcal{V}' accepts with probability $\geq c$

$x \notin L \Rightarrow$

- For any π , \mathcal{V} rejects with probability $\geq \epsilon$
- At least *epsilon* fraction of outcomes $\mathcal{V}'_{x,r}(\pi)$ are rejections

2-query PCP

$x \in L \Rightarrow$

- $\exists \pi$ that \mathcal{V} accepts with probability $\geq c$
- $\langle \text{Perm}(\pi), \pi \rangle$ is a proof that \mathcal{V}' accepts with probability $\geq c$

$x \notin L \Rightarrow$

- For any π , \mathcal{V} rejects with probability $\geq \epsilon$
- At least *epsilon* fraction of outcomes $\mathcal{V}'_{x,r}(\pi)$ are rejections
- In any proof π_1, π_2 , π_1 need not necessarily be $\text{Perm}(\pi_2)$

2-query PCP

$x \in L \Rightarrow$

- $\exists \pi$ that \mathcal{V} accepts with probability $\geq c$
- $\langle \text{Perm}(\pi), \pi \rangle$ is a proof that \mathcal{V}' accepts with probability $\geq c$

$x \notin L \Rightarrow$

- For any π , \mathcal{V} rejects with probability $\geq \epsilon$
- At least *epsilon* fraction of outcomes $\mathcal{V}'_{x,r}(\pi)$ are rejections
- In any proof π_1, π_2 , π_1 need not necessarily be $\text{Perm}(\pi_2)$
- If $\pi_1(i_1, \dots, i_q) = \pi_2(i_1, \dots, i_q)$, for *epsilon* fraction of the runs, it is rejected

2-query PCP

$x \in L \Rightarrow$

- $\exists \pi$ that \mathcal{V} accepts with probability $\geq c$
- $\langle \text{Perm}(\pi), \pi \rangle$ is a proof that \mathcal{V}' accepts with probability $\geq c$

$x \notin L \Rightarrow$

- For any π , \mathcal{V} rejects with probability $\geq \epsilon$
- At least *epsilon* fraction of outcomes $\mathcal{V}'_{x,r}(\pi)$ are rejections
- In any proof π_1, π_2 , π_1 need not necessarily be $\text{Perm}(\pi_2)$
- If $\pi_1(i_1, \dots, i_q) = \pi_2(i_1, \dots, i_q)$, for *epsilon* fraction of the runs, it is rejected
- Otherwise, $\pi_1(i_1, \dots, i_q)$ differs from $\pi_2(i_1, \dots, i_q)$ in at least one position, and the probability of \mathcal{V}' rejecting it is $1/q$

2-query PCP

$x \in L \Rightarrow$

- $\exists \pi$ that \mathcal{V} accepts with probability $\geq c$
- $\langle \text{Perm}(\pi), \pi \rangle$ is a proof that \mathcal{V}' accepts with probability $\geq c$

$x \notin L \Rightarrow$

- For any π , \mathcal{V} rejects with probability $\geq \epsilon$
- At least *epsilon* fraction of outcomes $\mathcal{V}'_{x,r}(\pi)$ are rejections
- In any proof π_1, π_2 , π_1 need not necessarily be $\text{Perm}(\pi_2)$
- If $\pi_1(i_1, \dots, i_q) = \pi_2(i_1, \dots, i_q)$, for *epsilon* fraction of the runs, it is rejected
- Otherwise, $\pi_1(i_1, \dots, i_q)$ differs from $\pi_2(i_1, \dots, i_q)$ in at least one position, and the probability of \mathcal{V}' rejecting it is $1/q$
- Hence, overall test rejects with probability at least ϵ/q

Summary

Thank you!